



Gene selection for cancer type classification

The purpose of our project is to work on a high-dimensional genomics data and find a relatively small number of genes to predict the cancer type of a given tumorous cells. This is known as Genes Selection for Cancer Classification and it is in line with many up-to-date problems of applied medicine.

Our dataset contained 1.032 cancerous cells and their knock-out probabilities, i.e. the probabilities of stopping the growth of tumor by inhibiting one of the ~ 17.000 genes. Each cell line was characterized by one of 10 possible cancer labels: Eye, Gastrointestinal, Gynecologic, Musculoskeletal, Neurological, Breast, Head-Neck, Hematologic, Genitourinary and, finally, Lung. We explored in details three Features Selection algorithms: Random Forests (RF) combined with Feature Importance, Lasso-SVM and Neural Networks (NN) combined with Olden Importance. We studied two binary classification problems (Blood-cancer vs. All, Lung-cancer vs. All) and the multiclass problem. Besides lung models, we achieved satisfying classification accuracies and we were able to select about 100-200 genes from the initial ~ 17.000 ones. Models fitted on such genes obtained classification accuracies ranging from 67% to 98%.

Therefore, it seems that classifying cancer type from an extremely small set of genes depends on the cancer type itself. These methodologies worked incredibly well on Blood cancer, as we reached almost 100% accuracy with the reduced classifier, while failed miserably on Lung cancer.

1 Introduction

Cancer is a complex disease characterized by the uncontrolled growth of abnormal cells anywhere in the body. These abnormal cells are extremely invasive and we usually identify them with the name of their original tissue (for instance, breast cancer, lung cancer, brain cancer, etc.). In recent years, medicine has made a great step forward in finding new and efficient therapies for different diseases and cancer is one of them. In particular, thanks to numerous advances in technology, collecting huge amount of data is no longer an issue, so that one can exploit such information to define personalized treatments for patients. In this regards, the DepMap project¹ and, in particular, the Achilles project² aim to use genome-wide screens to collect data regarding mutations of cancerous cells, identify essential genes and report vulnerabilities across hundreds of human cancers.

Many researches are currently using DepMap datasets to identify relatively small sets of genes which are responsible of cancers growth³. This procedure is often driven by medical knowledge, which we do not possess, together with some rough measures of importance. Being Maths student, we instead base our research on statistical models and on the hypothesis that, if a given classifier is able to distinguish different types of cancer, then the most relevant genes are the most important features for that classifier (the meaning of "important features" will be clarified later). Of course, selecting few truly significant genes has outstanding implications in the medical field: building faster diagnosis tools and synthesizing less toxic drugs are only two examples.

¹DepMap Portal: <https://depmap.org/portal/>

²Achilles Project: <https://depmap.org/portal/achilles/>

³Background material: <https://depmap.org/portal/publications/>

2 Dataset

We used two public datasets from the DepMap Public 21Q3 database, released on August 2021⁴:

- D1 *CRISPR_gene_dependency.csv*, which contains 1.032 cancer cells and their 17.393 gene scoring results;
- D2 *sample_info.csv*, which contains cell lines information, such as primary disease and sample collection site.

Data were collected from real patients and successively processed, so that element (i, j) of this (1.032×17.393) -data frame is the probability that "knocking out gene j has a real depletion effect on the i -th cell". Before proceeding with our analysis, we removed missing values, which affected only 10 rows coming from different tumours. Regarding weird observations, we found 2 "Non-Cancerous" and 6 "Engineered" cells. The first can be reasonably discarded, whereas the latter requires a little care. Engineered cells are synthetically modified samples in lab and, here, they are manly associated to the Eye sample collection site. We decided to keep them and associate them to the cancer corresponding to their site.

We grouped the various cancer types in 10 classes according to common medical knowledge⁵ and we obtained classes as reported in Figure 1. "Eye" was the smallest one as there were only 16 observations, 5 of which labelled as "Engineered". On the other hand, "Gastrointestinal" was the largest group and it comprehended 7 types of cancer, making this group quite heterogeneous.

We investigated two binary classification problems, Blood vs Rest and Lung vs Rest, and the multiclass problem. We chose "Lung" because of the nature of such a class: it was the most numerous group composed only by Lung cancer samples. The choice of "Blood" was instead

driven by some underlying biological knowledge: Blood cancer is quite different from other tumours because

- blood is in the whole body, and so the cancer is, too;
- Leukemia, Lymphom and Myeloma are the main kinds of cancer but they all affect white blood cells;
- not all blood cancers require a treatment, just periodical monitoring.

3 Methods

Let us briefly illustrate the algorithms we used. Our methodology was characterized by three steps:

1. fit the model using all the features;
2. identify the most important variables based on some measure of importance;

⁴Download dataset from DepMap: <https://depmap.org/portal/download/>

⁵Cancer types grouped by body location: <https://www.cancer.gov/types/by-body-location>




Figure 1: Cancer classes

3. use these genes to fit a reduced version of the classifier and find out its performance.

Clearly, each procedure involved fitting a model twice: the all-features version and the reduced one. We were thus forced to split the dataset into two further chunks. This was crucial to ensure independence of the two models and remove any sort of correlation.

3.1 Random Forest

We started by fitting Random Forest (RF) models as they frequently performs well on imbalanced and correlated high-dimensional data. Once the model had been fitted, we used Variable Importance of RFs to identify the most important features. This measure is calculated in three steps. First, prediction accuracy are measured on the out-of-bag samples. Then, the values of the variable are randomly shuffled, keeping all other variables the same. Finally, the decrease in prediction accuracy on the shuffled data is measured and the mean decrease in accuracy across all trees is reported.

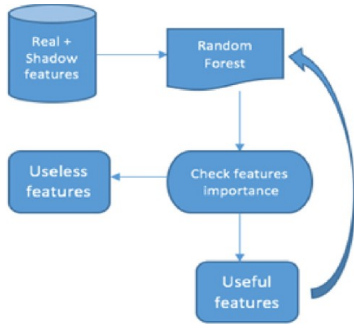


Figure 2: Boruta algorithm

Intuitively, the random shuffling means that, on average, the shuffled variable has no predictive power.

Hence, Variable Importance measures how much accuracy decreases because of removing a variable. Here, we used it to select the most important feature in two different ways:

- *Cross-validation*: we performed a 5-fold Cross-validation on the model and selected the top most important features of each model;
- *Boruta algorithm*⁶: Boruta repeatedly measures feature importance then carries out statistical tests to screen out the features which are irrelevant. The procedure terminates when all features are either decisively relevant or decisively irrelevant. Simple scheme is drawn in Figure 2.

3.2 SVM-Lasso

Support vector machines (SVM) are based on the idea of finding a hyperplane that best separate classes. Here, we combined this method with the classical Lasso penalty, so that the objective function to be minimized was:

$$\frac{1}{n} \sum_{i=1}^n \text{hingeLoss}(y_i(x_i w + t)) + \lambda \|w\|_1 \quad \text{where} \quad \text{hingeLoss}(z) = \max\{0, 1 - z\}$$

The parameter λ has been chosen via cross-validation. Thanks to Lasso penalty, we obtained parsimony in predictors: some w_i are shrunk all the way to zero, whereas the others identified important features.

3.3 Neural Networks

Neural Networks (NN) are efficient models to capture non-linear relationship between predictors and target variables. In this context, we trained NN with two hidden layers of width 400 and 300 respectively and we chose *ReLU* as activation function for the hidden layer, *sigmoid* and *softmax* for the output layer of, respectively, binary and multiclass classifications.

Given that the binary problems were a little unbalanced, we tried both the usual *Cross Entropy* loss function and the *Focal Loss*, defined as

$$FL(z) = \alpha \cdot (1 - z)^\gamma \log z, \quad \text{with } z \in [0, 1] \text{ and } \alpha, \gamma \geq 0$$

⁶Boruta algorithm: https://www.researchgate.net/publication/220443685_Boruta_-_A_System_for_Feature_Selection

Note that *Focal Loss* can be extended for dealing with multiclass classification tasks⁷.

Once our NNs were fitted, we ranked variables according to the Olden Importance measure⁸, selected the first ones and trained a reduced version of the classifier on them.

4 Results

4.1 Binary classifications

Before starting our Binary classifications on Blood and Lung cancer, we ran a Principal Component Analysis to gain some valuable insights. We noticed a surprising result: even if observations formed a cloud of points, Blood cancer cells were mainly concentrated in just one part of the 3-dimensional plot. The same did not happen for Lung cancer observations.

We then moved to RF classifiers. Unfortunately, we could not build a proper model for Lung vs Rest. Even after a tuning on tree parameters, our RF was not able to detect the minority class. Thanks to SMOTE, we adjusted the frequency of those observations from 11% to 50% and we refit another RF. Slightly better: at least one third of Lung cancer cells was correctly predicted. We also focused on pruning subtrees of RF. In particular, we looked for an optimal complexity parameter *ccp_alpha* for Minimal Cost-Complexity Pruning by means of a Randomized search. However, this new RF classifier built with such hyperparameter had reached the same accuracy of the one fitted on balanced data. We reported this accuracy in Table 1. Since all our models were not good enough in classifying Lung cancer, we retained that selecting the most important features was pointless and incorrect in terms of real applications.

On the other hand, we obtained great results in Blood vs Rest classification. Even though Blood cancer observations were only the 11% total ones, here we did not need any adjustments for the minority class. Indeed, we just tuned trees parameters and took a correction on the class weights to reach 98% of mean accuracy. For completeness, we fitted also a RF with Cost-Complexity Pruning and, again, we find the same accuracy.

Then, we removed unnecessary features by performing a 5-fold Cross-validation. Thanks to Variable Importance values, we could indeed select the most relevant genes from each of the 5 models and took only the ones in common. A more rigorous way to pursue variable selection is by applying the Boruta algorithm. As shown in Table 2, Boruta individuated a higher number of important features than our manual method and, in particular, they agreed only on 79 genes. Hence, we fit a RF by considering only these latter genes. As mentioned above, we used another dataset and now Blood cancer instances accounted by 9% of the total ones. Besides the correction of class weights, no further parameters were specified. Nevertheless, the new RF classifier predicted correctly 6 tumour cells out of 7.

As for the implementation of NNs, we firstly fit a single

Finally, SVM.

Table 1: Average recall of the models fitted with all the features

Task	RF	SVM-Lasso	NN
Blood	50	837	970
Lung	47	877	230
Multiclass	45	300	556

4.2 Multiclass classification

When working with the multiclass problem we decided to remove the group "Eye" because it was a small heterogeneous group that even after tuning for RF and NN held the worst accuracy scores. "Eye" was a cluster

⁷Focal Loss: <https://arxiv.org/pdf/1708.02002.pdf>

⁸Olden Importance: https://depts.washington.edu/oldenlab/wordpress/wp-content/uploads/2013/03/EcologicalModelling_2004.pdf

Table 2: Selected variables

Task	RF	Boruta	SVM-Lasso	NN
Blood	50	837	23	970
Lung	47	877	44	230
Multiclass	45	44	300	556

Table 3: Average recall of reduced models

Task	RF	SVM-Lasso	NN
Blood	50	837	970
Multiclass	45	44	300

of different types of cancer: this didn't create problems when we were using a One vs All approach but with multiclass the effects of doing the cluster ourselves was more evident. So after removing the group "Eye", we did the same as before and split the dataset in two.

Random Forests

We considered two different random forests, the first with the same parameters as the one used in Blood vs All random forest. Looking at the confusion matrix we can see that most of the misclassification errors are data points classified as Gastrointestinal. Gastrointestinal was the biggest cluster we consider with a lot of different cancers so the genes (i.e. the features) important can be different within the cluster. To limit the weight of Gastrointestinal we decided to tune by hand the class weights. In this way we solve two problems: we reduced the importance of the group Gastrointestinal and we achieved at least one correct classification in each class. At this point we used cross validation to select five models and we extracted the most important features by looking at their relative importance in every model subjected to their presence in every model. We set a threshold to obtain (~ 1700) features and with this we created a reduced model that we fitted using the validation set we had created at the beginning.

SVM-Lasso

Neural Networks

We first fitted a toy NN with only one hidden layer to see how the net worked. We then fitted three different NN: the first increasing by one the number of hidden layers, the second changing the loss function with the Focal loss and the third using again the Focal loss, but changing the alpha parameter w.r.t. the class percentage in the set. The three models held similar result with the second one being slightly more accurate. Once these NNs were fitted, we ranked the variable according to the Olden Importance measure, selected the first 1700 and used them to construct a reduced version of the classifier on the validation set. The increment in the number of features used for fitting the reduced model w.r.t. the binary NN model was because we are now trying to separate nine different classes so we need much more information than before.

5 Overall Results

We approached the multiclass problem using three different techniques: RF, NN, SVM-Lasso and all three methods gave us good results. We can't say the same for the reduced models: while NN maintain an overall good performance using the same models, the RF need to be retuned to reach a similar performance.

6 Conclusion and future works

- reiterate key findings
- limitations: data high correlated, no medical knowledge, no meaning at selected genes
- suggest for the future: work with a med student, try other classes or methods, improve lung in some way